

Numerical Simulation and Scientific Computing I

Lecture 2: Serial Code Optimization



Clemens Etl, Paul Manstetten, and
Josef Weinbub



Institute for Microelectronics
TU Wien

nssc@iue.tuwien.ac.at

Quiz – Question 1

- What is the minimum sum of exercise points you need to be eligible for the exam?

Quiz – Question 1

- What is the minimum sum of exercise points you need to be eligible for the exam?
 - Answer: at least 21 (of $30=3 \times 10$)

Quiz – Question 2

- Read the specification of this CPU

https://en.wikichip.org/wiki/intel/core_i9/i9-13900k

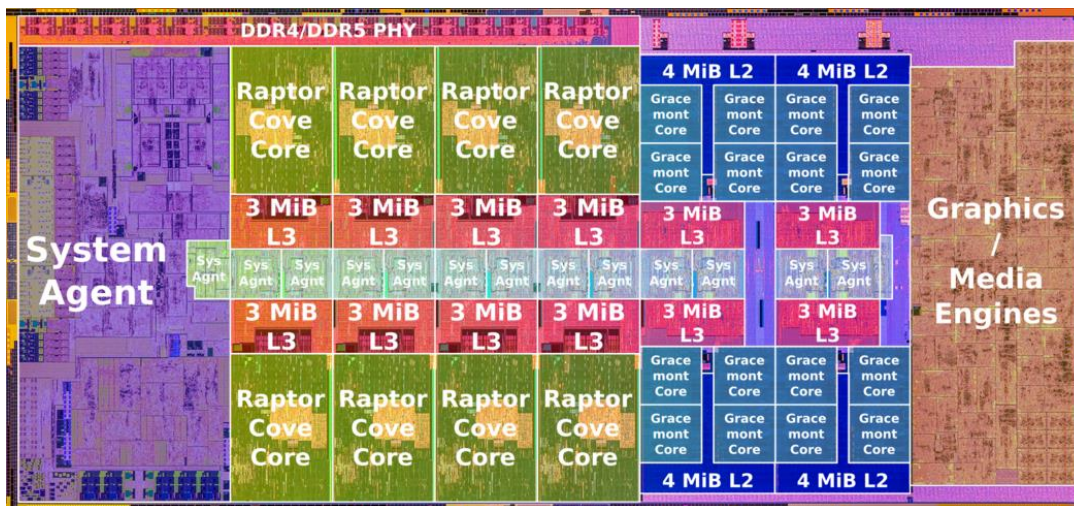
and classify it according to:

- Computer category
- Supported forms of architectural parallelism
- Flynn's taxonomy
- Instruction set architecture
- Memory hierarchy (cache levels)
- Type of multithreading
- Type of multicore architecture

Quiz – Question 2

Intel i9-13900k

- Brand new 2022, x86, 64bit Desktop CPU
 - 8 “big cores” (3-5.5GHZ), L1 I32K/D48K, L2 8x2M, L3 8x3M
 - 16 “small cores” (2-4.4GHZ), L1 I64K/D32K, L2 4x4, L3 4x3M
- Flynn’s Taxonomy
 - SISD: yes (single processor, ILP)
 - SIMD: yes (vector extensions: AVX2)
 - MIMD: yes (8 + 16 cores)
- Multithreading: SMT and thread-level
- Tightly-coupled MIMD: two-way SMP



Quiz – Question 3

- Caches are very expensive. Why do we need them at all?
 - Situation without caches (or very small caches)?

Quiz – Question 4

- What are simple ways you can use a compiler to make your program run faster?

Quiz – Question 5

- Why are SIMD extensions (e.g. AVX-512) not considered instruction-level parallelism?

Outline

Optimization stages

Performance modeling for a
“cache-based stored-program microprocessor”

The Optimizing Compiler

“It is easier
to make a correct program fast than
to make a fast program correct.”

Outline: Optimization Stages

- Basic Choices
 - General purpose vs. specialized hardware
 - Language / language “infrastructure”
 - Algorithms
 - Data structures
- Compile time
 - Static compilation ahead of time (AOT)
- Link time
 - Link time optimization (LTO) across module boundaries
- Run time
 - Using run time statistics for profile guided optimization (PGO)
- Summary

Compute Platform

- Choice of suitable and available compute platform
 - General purpose CPU: Intel, AMD, ARM, ...
 - General purpose GPU: Nvidia, AMD, ...
 - Special hardware, e.g., Google's TPUs, Bitcoin Miner ASICs, ...
- Example for specified theoretical max. performance metrics

*table data from 2019	CPU W-2195	GPU Tesla T4	TPU Google v2	ASIC Antminer S9
Floating point operations	1710 GFlops/s (FP64)	8100 Gflops/s (FP32)	180 000 GFlops/s (FP16)	14 000 GHashes/s sha256(sha256(80byte))
Memory bandwidth	85.3 GByte/s	320 GByte/s	--	--
TDP	140W	70W	75W	1300W
Price	~3000USD	~3000USD	--	~3000USD

Programming Language

- Choice of programming language
 - Experience
 - Features of the language
 - Available runtimes/compilers/libraries
 - Interfacing to required resources or other language/runtimes
 - Portability
 - Performance requirements

Language	Implementations/ Runtimes/Compilers	Recent Versions	Comment
Python	CPython, PyPy	(python2), python3	interpreted, native extensions
JavaScript	V8,JSC,JS	ECMAScript15(“ES6”)	interpreted, JIT comp. for optimization
C++	libc++ with clang, libstdc++ with gcc	C++11/14,C++17	statically typed native

Algorithms / Data Structures

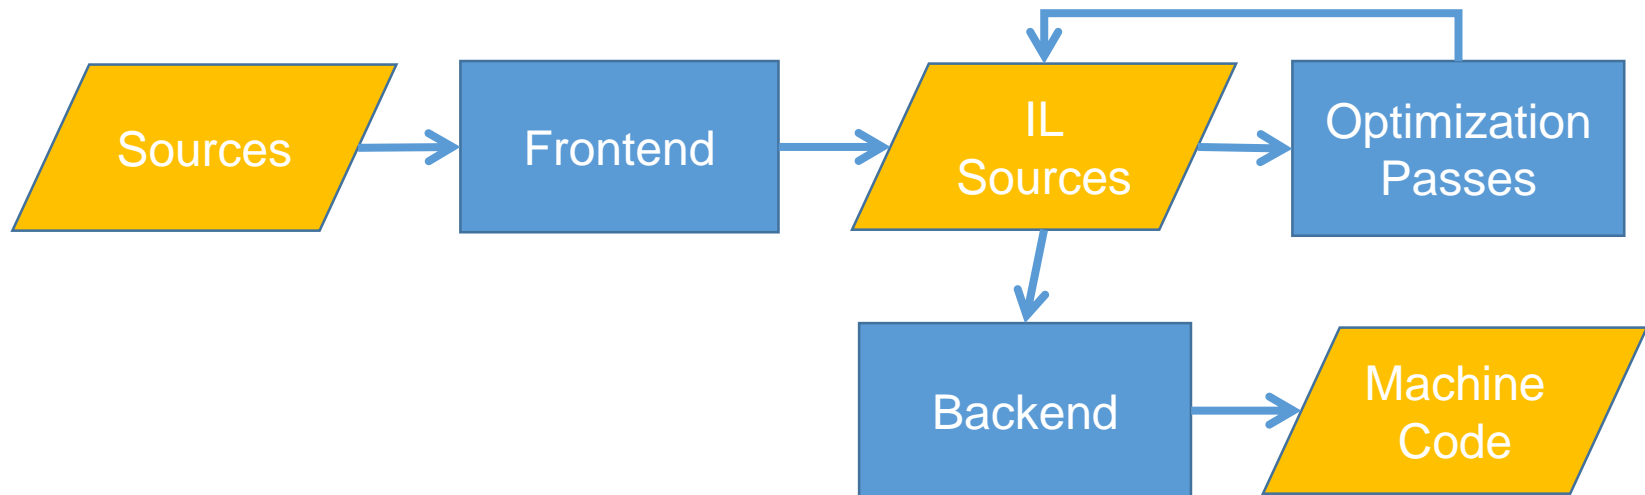
- Choice of suitable algorithm for each computational task
 - Consider the algorithmic complexity w.r.t. problem size
 - Consider the underlying data structure w.r.t.
 - memory footprint for anticipated problem sizes
 - anticipated access patterns
 - Tradeoff with implementation effort
- Practically achieved performance greatly influenced by ...
 - the implementation on a particular platform,
 - the considered problem size,
 - run time access patterns,
 - which is all “hidden” in a prefactor “ $K * O(\dots)$ ”.

Examples of complexity guarantees for containers in the C++ standard library

Container	Insert	Read	Erase
<code>std::vector<T></code>	$O(1)$, $O(N)$	$O(1)$	$O(1)$, $O(N)$
<code>std::map<T1,T2></code>	$O(\log N)$	$O(\log N)$	$O(\log N)$
<code>std::unordered_map<T1,T2></code>	$O(1)$, $O(N)$	$O(1)$, $O(n)$	$O(1)$, $O(n)$

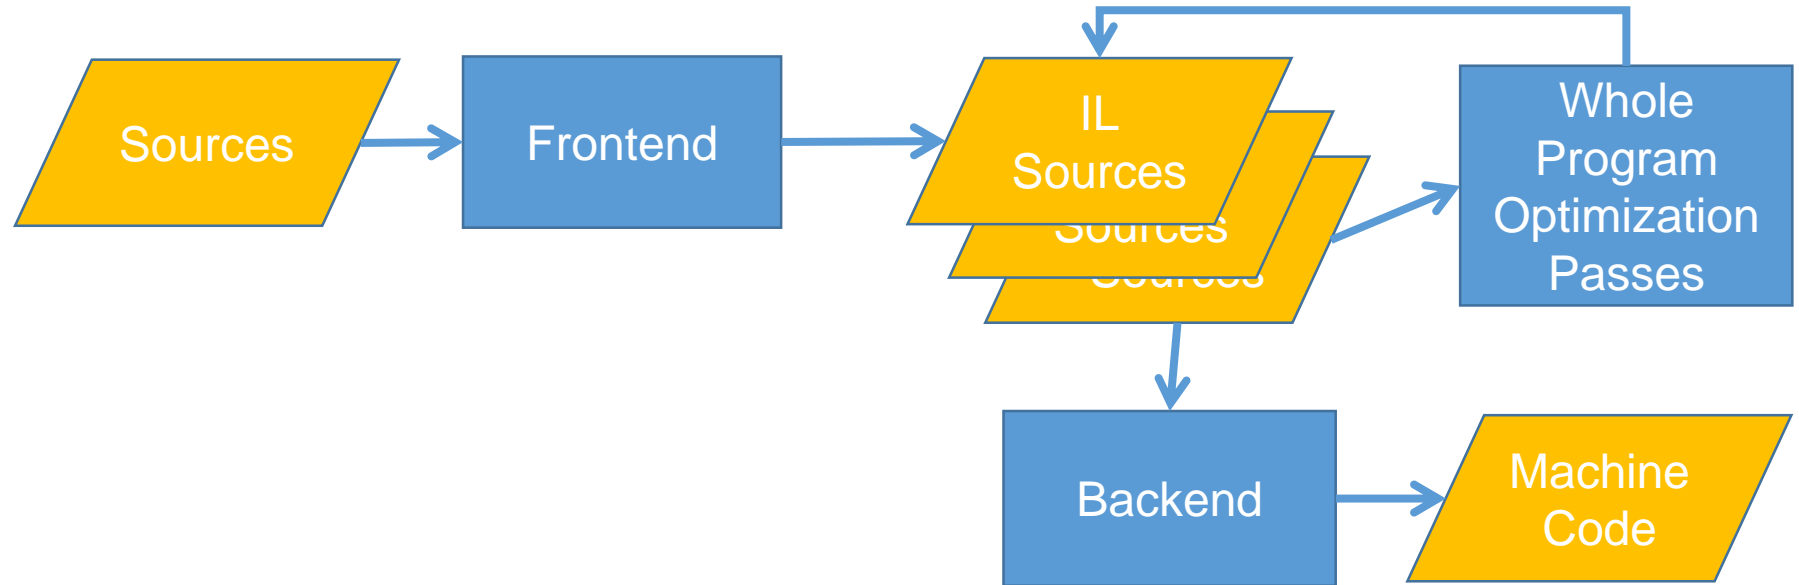
Optimization At Compile Time (C++)

- Optimization flags control static code optimization during AOT compilation
 - Various optimizations for loops (unroll, interchange, vectorize), function calls (inlining), reordering of independent instructions
- Optimization is affected by target architecture
 - # of Registers, special registers (SIMD), # functional units, pipelining capabilities, cache hierarchy/sizes
- Optimization scope is limited to individual compilation units



Optimization At Link Time (C++)

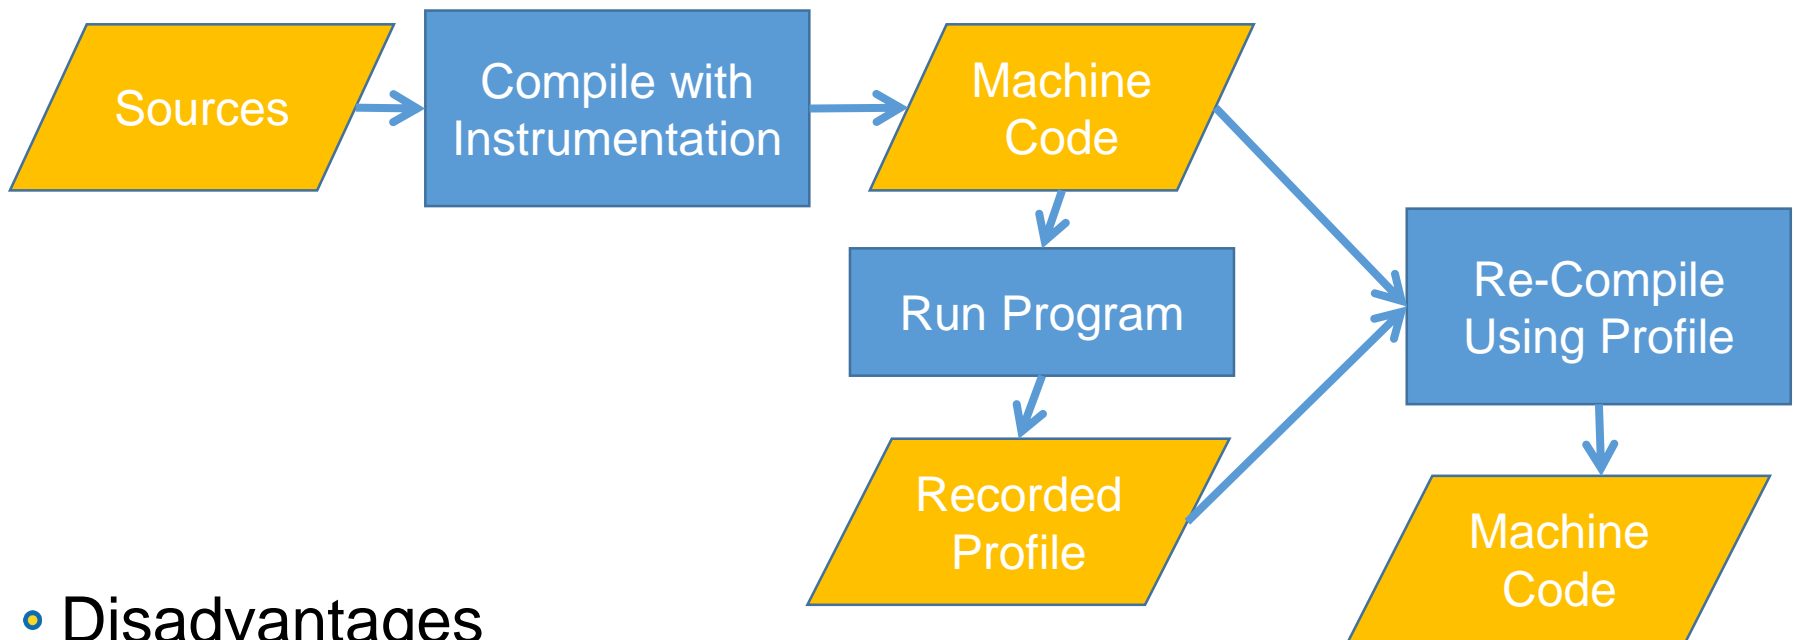
- Whole program optimization (“-flto”)
 - Function inlining across module boundaries, dead code elimination



- Disadvantages (as with larger “header-only” projects)
 - Long compilation times
 - Requires a lot of memory, hard to parallelize

Optimization At Run Time (C++)

- Profile guided optimization (PGO)
 - Optimization decisions based on run time profile
 - Also attractive for JIT compilation used by interpreters: use run time statistics to statically compile “hot” regions of the code



- Disadvantages
 - Optimized for input used to generate profiles: unpredictable performance for different inputs
 - Multiple compiles necessary

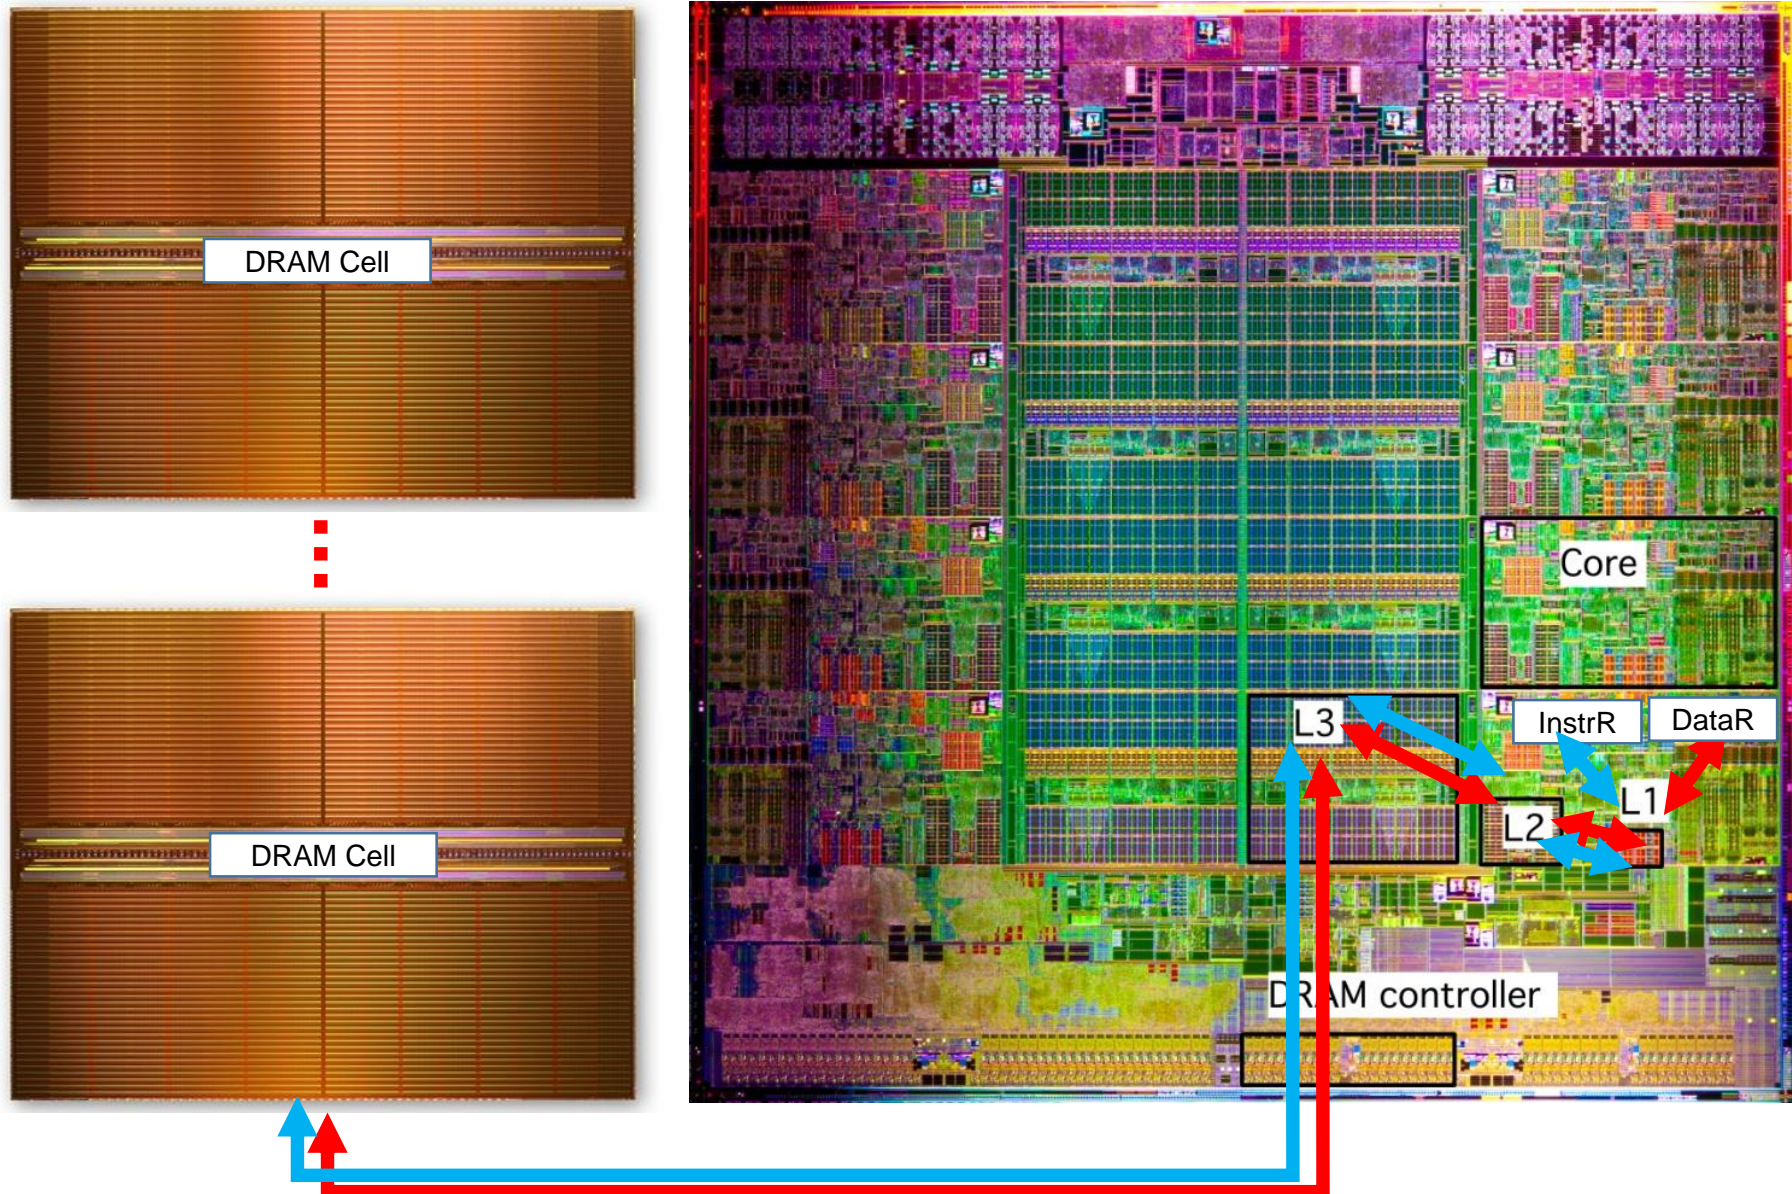
Summary: Optimization Stages

- Optimization Stages
 - Platform / Data Structures / Algorithm / Language
 - Often this choice is constrained
 - Compile Time (AOT)
 - Static optimizations on module level
 - Link Time (LTO)
 - Static optimizations between modules
 - Run Time (PGO)
 - Optimization using run time statistics to guide opt. decisions
- Python (interpreted)
 - For performance: native C/C++ extensions (e.g., numpy)
- JavaScript (interpreted)
 - For performance: JIT= \sim PGO, precompiled/typed extensions
- C++ (compiled)
 - For performance: AOT compiled, LTO, PGO

Outline: Performance Modeling

- Considered Architecture:
 - “Cache-based stored-program” x86 microprocessor
- Bandwidth-based modeling
- Data sizes
- Optimization potential estimation
- Effective bandwidth benchmarking: “Vector Triad”

Cache-Based Stored-Program x86 MC



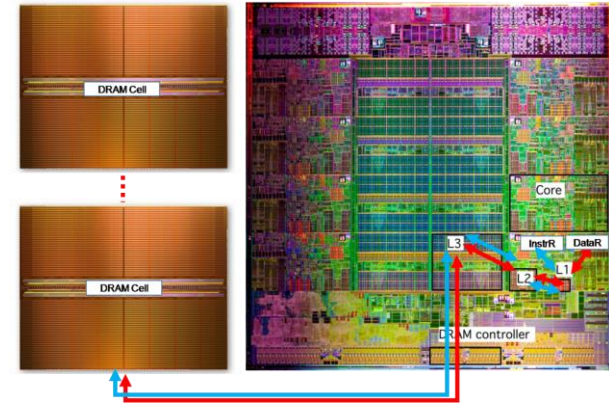
Bandwidth-Based Modeling

- Readily accessible for most systems:
 - Peak performance = maximum theoretical possible Flops/s
 - Including full SIMD utilization, including “FMA”
 - Underlying floating point representation depends on platform
 - Memory bandwidth = maximum theoretical main memory bandwidth
 - A single core might not be able to utilize this maximum
- Machine Balance

$$B_{machine} = \frac{\text{memory bandwidth [Bytes / s]}}{\text{peak performance [Flops / s]}} = \frac{\text{Bytes}}{\text{Flops}}$$

- Code Balance

$$B_{code} = \frac{\text{data traffic [Bytes]}}{\text{floating p. operations [Flops]}} = \frac{\text{Bytes}}{\text{Flops}}$$



Bandwidth-Based Modeling

- Machine Balance

$$B_{machine} = \frac{\text{memory bandwidth [Bytes / s]}}{\text{peak performance [Flops / s]}} = \frac{\text{Bytes}}{\text{Flops}}$$

- i7-4790K (2014), all cores


$$B_{machine} = \frac{25.6 [\text{GBytes / s}]}{256 [\text{GFlops / s}]} = 0.1 \frac{\text{Bytes}}{\text{Flops}}$$

- i7-940 (2008), all cores

$$B_{machine} = \frac{25.6 [\text{GBytes / s}]}{47 [\text{GFlops / s}]} = 0.54 \frac{\text{Bytes}}{\text{Flops}}$$

- Tesla T4 (2018), single precision FP

$$B_{machine} = \frac{85.3 [\text{GBytes / s}]}{1710 [\text{GFlops / s}]} = 0.05 \frac{\text{Bytes}}{\text{Flops}}$$



Bandwidth-Based Modeling

- Example

- i7-4790K (2014), all cores

$$B_{machine} = \frac{25.6 [GBytes / s]}{256 [GFlops / s]} = 0.1 \frac{Bytes}{Flops} \quad I_{machine} = 10 \frac{Flops}{Bytes}$$

- Benchmark code: “Vector Addition” (Large N)

$$B_{code} = \frac{3 \cdot 8 \cdot N [Bytes]}{1 \cdot N [Flops]} = 24 \frac{Bytes}{Flops} \quad I_{code} = 0.04 \frac{Flops}{Bytes}$$

- Expected fraction of theoretical peak performance

$$\min \left(1, \frac{B_{machine}}{B_{code}} \right) = \min \left(1, \frac{I_{code}}{I_{machine}} \right) = \min \left(1, \frac{0.04}{10} \right) = 0.004$$

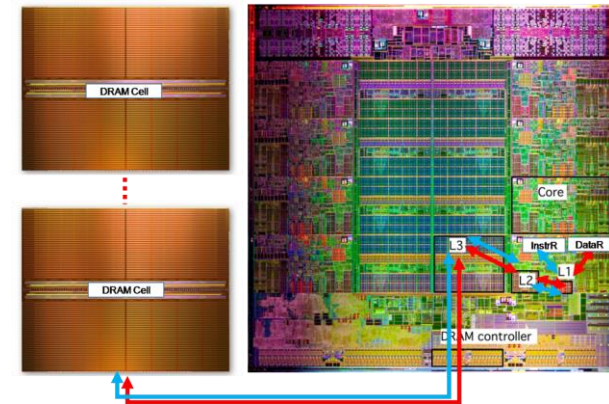
- Estimate peak performance and memory bandwidth (single core) more realistic (no SIMD & single core)

$$P_{peak(fullSIMD,allcores)} = 4[GHz] \cdot 4[Cores] \cdot 2[SIMD_{units}] \cdot 4[SIMD_{width}] \cdot 2[Flops(FMA)] = 256[GFlops / s]$$

$$P_{peak(noSIMD,onecore)} = 4[GHz] \cdot 1[Cores] \cdot 2[SIMD_{units}] \cdot 1[SIMD_{width}] \cdot 1[Flop] = 8[GFlops / s]$$

$$I_{machine} = \frac{256 / 32 [GFlops / s]}{25.6 / 2 [GBytes / s]} = 0.6 \frac{Flops}{Bytes}$$

$$\min \left(1, \frac{I_{code}}{I_{machine}} \right) = \min \left(1, \frac{0.04}{0.6} \right) = 0.07$$



Vector Addition:
do i=1,N
 A[i] = B[i] + C[i]
enddo

Data Sizes

- RISC Instructions


- Size: 32bit = 4bytes
- Types: data handling, arithmetic/logic, control flow

- Integer types

- From 8bit to 64bit, signed and unsigned
- Unsigned range 8bit: $2^{**}8=256$
- Unsigned range 64bit: $2^{**}64=18,446,744,073,709,551,615$

- Floating point types

- FP64, 64bit, double precision: “standard for scientific computing”
 - $1 - 1/3 = 1.3333333333333333$
- FP32, 32bit, single precision: “standard for computer graphics”
 - $1 - 1/3 = 1.3333334$
- FP16, 16bit, half precision: supported on TPUs/recent GPGPUs
 - $1 - 1/3 = 1.333$



Optimization Potential

- Scalar vector product

$$B_{code} = \frac{O(N) [Bytes]}{O(N) [Flops]} = \frac{3 \cdot 8 \cdot N [Bytes]}{2 \cdot N [Flops]} = 12 \frac{Bytes}{Flops}$$

```
Scalar product:  
do i=1,N  
  A[i] += B[i] * C[i]  
enddo
```

- Matrix addition

$$B_{code} = \frac{O(N^2) [Bytes]}{O(N^2) [Flops]} = \frac{3 \cdot 8 \cdot N^2 [Bytes]}{1 \cdot N^2 [Flops]} = 24 \frac{Bytes}{Flops}$$

```
Matrix addition:  
do i=1,N  
  do j=1,N  
    A[i,j] = B[i,j] + C[i,j]  
  enddo  
enddo
```


- Matrix multiplication

$$B_{code} = \frac{O(N^2) [Bytes]}{O(N^3) [Flops]} = \frac{3 \cdot 8 \cdot N^2 [Bytes]}{2 \cdot N^3 [Flops]} = 12 / N \frac{Bytes}{Flops}$$

```
Matrix Multiplication:  
do i=1,N  
  do j=1,N  
    do k=1,N  
      A[i,j] += B[k,j] * C[i,k]  
    enddo  
  enddo  
enddo
```

Example: “Vector Triad”

- Smallest data unit: cache line
 - On x86: 64 Bytes (= 8 FP64, =16 FP32)
- Hierarchy of bandwidths/latencies
 - L1/L2/L3/memory
- Performance bounds for “Vector Triad”
 - Large N, non-sequential access, memory bound, latency hidden
 - Large N, sequential access, memory bound, latency hidden
 - Infinite bandwidth, computationally bound



$$runtime = \frac{N \cdot 4 \cdot 64 [Bytes]}{bandwidth [Bytes / s]}$$

$$runtime = \frac{N \cdot 4 \cdot 8 [Bytes]}{bandwidth [Bytes / s]}$$

$$runtime = \frac{N \cdot 2 [Flops]}{peak\ performance [Flops / s]}$$

“Vector Triad”:

do i=1,N

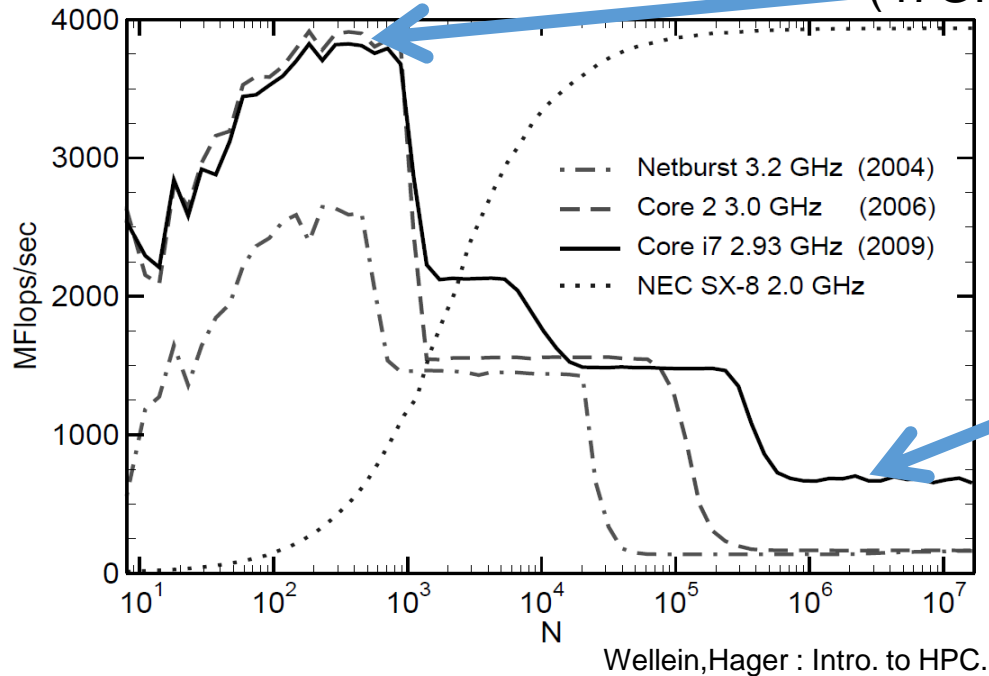
A[i] = B[i] + C[i] * D[i]

enddo

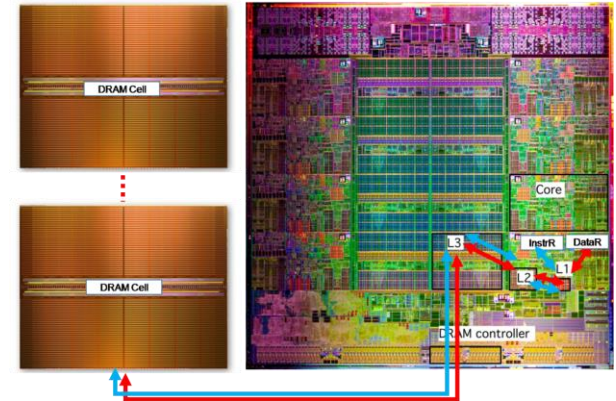
$$B_{triad} = \frac{O(N) [Bytes]}{O(N) [Flops]} = \frac{4 \cdot 8 \cdot N [Bytes]}{2 \cdot N [Flops]} = 16 \frac{Bytes}{Flops}$$

Example: “Vector Triad”

- Results for “Vector Triad”



Peak perf. ?
(47GFlops/s)



Max. bandwidth ?
(25.6GB/s)

“Vector Triad”:

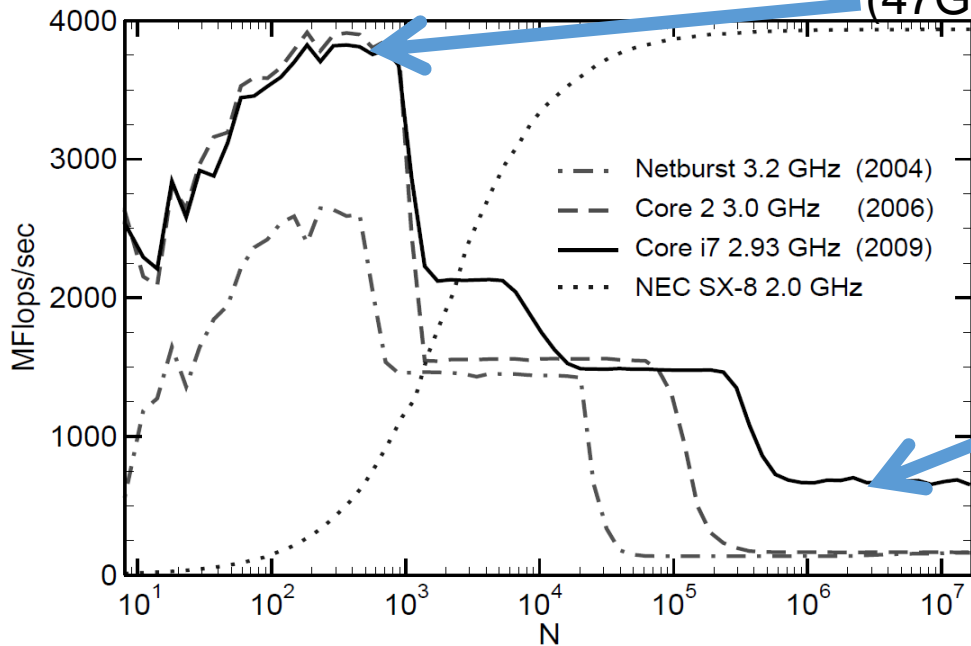
```
do i=1,N  
  A[i] = B[i] + C[i] * D[i]  
enddo
```

- Interpretation

- Very small N: pipelines not yet efficient
- Small N: data contained in L1 cache, cache bandwidth bound
- Medium N: steps manifest cache bandwidth levels
- Large N: bound by streaming bandwidth from memory
- Vector processor: complementary results

Example: “Vector Triad”


- Results for “Vector Triad”



Wellein, Hager : Intro. to HPC.

Peak perf. ?
(47GFlops/s)

Max. bandwidth ?
(25.6GB/s)



“Vector Triad”:

```
do i=1,N
  A[i] = B[i] + C[i] * D[i]
enddo
```

- N~500: 3.8 Gflops/s

- Effective L1 bandwidth: $3.8 [GFlops / s] \cdot B_{triad} = 61 [GBytes / s]$

- N=Large: 0.75Gflops/s

- Effective memory bandwidth: $0.75 [GFlops / s] \cdot B_{triad} = 12 [GBytes / s]$
(for a single core)

Summary: Performance Modeling

- Smallest data unit: cache line (64bytes)
- Data sizes
 - Instructions (4Bytes)
 - Integral types (1-8Bytes)
 - Floating point types (2-8Bytes)
- Bandwidth-based modeling
 - Model for bandwidth limited code (disk, memory, caches)
 - Machine balance / Code balance
- Optimization potential estimation
 - Ratio of data traffic vs. floating point operations
- Example Code: “Vector Triad”
 - Benchmarking effective bandwidths of a system (single core)

The Optimizing Compiler

- C++ compiler flags for performance optimization
 - “-O0”: no opt., most debugable code
 - “-O2”: enabling most optimization
 - “-O3”: more optimization (longer compile time, pot. larger code)
 - “-O3 -ffast-math”: allows optimizations violating IEEE FP standards
 - “-flto”: enable link time optimization
 - “-march=native”: select compiling machine arch. as target
- Some important automatic compile time optimizations
 - Local optimization
 - Inlining
 - Vectorization
 - Loop optimizations
- Goals of compile time optimizations for performance
 - Optimally utilize target capabilities: Register, pipelines, functional units, caches
 - Avoid performance decrease for any possible run time input

Quiz

- Q1: How to deal with “turbo frequencies” when estimating the peak performance for a CPU?
- Q2: Why did Intel drop AVX-512 from the P-Cores (e.g. i9-13900k)?
- Q3: For very large N , is the performance of a dense matrix-matrix multiplication memory bound or computationally bound?

Quiz (for next Lecture)

- Q4: What are options to discretize the first and second derivatives of a one-dimensional function using discretization points of distance h ?
- Q5: Using the trapezoidal rule to integrate

$$A = \int_{-\pi/2}^{+\pi/2} \cos(x) dx$$

does the approximation overestimate or underestimate the integral A ?

Next Lecture: C++ Optimization, Finite Difference Method