

Numerical Simulation and Scientific Computing I

Lecture 6: Numerical Linear Algebra II – Krylov Methods



Clemens Etl, Paul Manstetten, and Josef
Weinbub



Institute for Microelectronics
TU Wien

nssc@iue.tuwien.ac.at

Quiz – Question 1

- Imagine you track the balance of a bank account using a single precision floating point number representing EUR. Starting with balance 0 EUR, each day 0.1EUR are transferred to the account, after how many days will the balance not increase anymore?

Poll 1

- What will be the order of magnitude of the EUR balance when it stops increasing (epsilon $\sim 10^{-7}$)?
 - A) 10^3
 - B) 10^6
 - C) 10^7
 - D) 10^9
 - E) 10^{12}

Quiz – Question 1

- Back-of-the-envelope calculation:
 - When will the scaled epsilon be on the order of 0.1?

$$\begin{aligned}x * 10^{-7} &\approx 0.1 \\ \rightarrow x &\approx 10^6\end{aligned}$$

Quiz – Question 1

```
{
    std::cout << "account balance" << std::endl;
    float f = 0.0f;
    double feps = std::numeric_limits<float>::epsilon();
    int days = 0;

    while (f != f + 0.1f) {
        f += 0.1f;
        ++days;
    }
    std::cout << "feps=" << feps << std::endl;
    std::cout << "balance=" << f << std::endl;
    std::cout << "days=" << days << std::endl;
    std::cout << "years=" << days / 365.0 << std::endl;
    std::cout << "decades=" << days / 365.0 / 10.0 << std::endl;
}
// output:
// arithmetics
// account balance
// feps=1.19209e-07
// balance=2.09715e+06
// days=18073720
// years=49517
// decades=4951.7
```

Quiz – Question 2

- What are potential advantages/disadvantages of using BLAS/LAPACK or Eigen?

Quiz – Question 2

- What are potential advantages/disadvantages of using BLAS/LAPACK or Eigen?
 - License (modified BSD vs MPL2)
 - Language (Fortran vs C++)
 - Speed (similar, but vs. implementation from scratch)
 - BLAS/LAPACK: more API than implementation

Quiz – Question 3

- How could you calculate an upper bound for the spectral radius for a given iteration matrix of the Jacobi method?

Quiz – Question 3

- Spectral radius (condition for convergence)
 - $\rho(D^{-1}(A - D)) < 1$
- Smaller than **any** norm
 - $\rho(X) \leq \|X\|$
 - Why? $\|X\| = \max_{x \neq 0} \frac{\|Xx\|}{\|x\|}$
- Choosing the **maximum norm**:
 - $\|D^{-1}(A - D)\|_{\infty} < 1 \rightarrow |a_{ii}| > \sum_{j \neq i} |a_{ij}|$
 - **Strict diagonal dominance** is a sufficient condition for convergence

Outline

- Motivation
- Data Structures
- Krylov Subspace
- Methods
 - GMRES
 - Conjugate Gradient (CG)
 - Bi-CG(STAB)
- Preconditioners

Goal

- Motivate the introduction of iterative methods
 - Key concept: sparsity
- Get a “feeling” of Krylov subspace methods
 - Including a short tour of popular choices
- Explore their limitations
- Some practical considerations
 - Implications for memory and data structures

What we will NOT cover

- Mathematical proofs
- Convergence analyses
- Eigenvalue problems
- Direct sparse solvers
- Multigrid
- Detail!

All that we care about

$$Ax = b$$

Take-home message

- Krylov methods can be useful when there are large, sparse matrices
 - We do not want to lose sparsity
 - We do not want to pay the $O(n^3)$ price
- Convergence can be hard
 - Choosing the correct method for the problem
 - Equal parts “art” and “science” – experience & experimentation are key!
- Preconditioners can significantly increase convergence rate
 - Also make your problem more stable

Main References

- Iterative Krylov Methods for Large Linear Systems
 - Author: Henk A. van der Vorst
 - eBook available:
https://catalogplus.tuwien.ac.at:443/UTW:UTW:TN_cambridge_s10_1017_CBO9780511615115
- Eigen documentation
 - https://eigen.tuxfamily.org/dox/group__TutorialSparse.html

Additional References

- Numerical Linear Algebra

- Authors: David Bau and Lloyd N. Trefethen
- https://catalogplus.tuwien.ac.at:443/UTW:UTW:UTW_alma2143358810003336

- Matrix Computations

- Authors: Gene H. Golub and Charles F. Van Loan
- https://catalogplus.tuwien.ac.at:443/UTW:UTW:UTW_alma2149765900003336

- Iterative Methods for Sparse Linear Systems

- Author: Yousef Saad
- https://catalogplus.tuwien.ac.at:443/UTW:UTW:UTW_alma2154052560003336

- Any examples of linear algebra libraries?

Libraries

- Eigen
 - “Eigen is a C++ template library for linear algebra: matrices, vectors, numerical solvers, and related algorithms”
- PETSc
 - “PETSc is a suite of data structures and routines for the scalable solution of scientific applications modeled by PDEs”
- SciPy
 - “SciPy is a Python-based ecosystem of open-source software for mathematics, science, and engineering”
- MTL4, Armadillo, Trilinos
 - Similar in spirit to Eigen
- SLEPc
 - “SLEPc, the Scalable Library for Eigenvalue Problem Computations, is a software library for the solution of large sparse eigenproblems on parallel computers”

Recap – LU Decomposition

- For a non-singular $n \times n$ matrix A , we want to solve:

$$Ax = b$$

- A decomposition exists such that:

$$PA = LU$$

- Where:

- P is a permutation matrix
- L is lower triangular
- U is upper triangular

- Therefore:

$$\begin{aligned} Ax = b &\Leftrightarrow Ly = Pb \\ Ux &= y \end{aligned}$$

Motivation – Computational Resources

- LU decomposition
 - $O(n^3)$ for dense matrices.
 - In more reasonable scenarios (banded matrix) - $O(n^{2\frac{1}{3}})$
- If the computational power increases 1000x – only a 10x increase in problem size

Sparsity

- Definitions:

- n is the matrix dimension $\rightarrow n^2$ total entries
- m is the number of non-zero entries

- Dense Matrix

- $m \sim n^2$
- Very common in scientific problems
- e.g. 5 chemical species in a reaction

1	4.6	0.7	0	3.8
9.3	1	8.5	3.7	0.7
2.3	6	2.7	7.6	5.9
1.2	0	0	4.8	9.4
7.1	10.9	5.6	1	1


- Sparse Matrix

- $m \ll n^2$
- Usually arise from discretization of integral or differential equations
- Rule of thumb: if n is very large, then it is probably an approximation to ∞

1	4.6	0	0	0
0	0	8.5	3.7	0
0	6	2.7	0	0
0	4.6	0	4.8	9.4
0	0	5.6	0	1

Motivation – Memory


- Source: Matrix Market – A repository for test data
 - <https://math.nist.gov/MatrixMarket/>
 - Example: Matrix MAN 5976 (Structural Engineering)



- Structure plot: color the non-zero entries
- **Bandwidth**: maximum separation between non-zero entries
- Is it necessary to store the whole matrix?


Poll 2

- What will happen to the sparsity of the matrix with the following structure plot after applying LU factorization?
 - A) Keeps the same sparsity
 - B) Loses some sparsity
 - C) Becomes completely dense



Poll 2

- What will happen to the sparsity of the matrix with the following structure plot after applying LU factorization?
 - A) Keeps the same sparsity
 - **B) Lose some sparsity**
(depending on the pivoting!)
 - C) Become completely dense



Quiz – Question 4

- Assume a large matrix has mostly zero entries, how to store it efficiently in terms of memory footprint?

Quiz – Question 4

- Assume a large matrix has mostly zero entries, how to store it efficiently in terms of memory footprint?

We don't want to store zeros!

Storing Sparse Matrices - COO

- COO: **Coordinate Format**

- Define 3 vectors of size m : V , IA and JA
- V contains all non-zero entries
- IA , JA contain the corresponding i and j indices
- Storage requirement: $3m$

$$\begin{bmatrix} 1 & 4.6 & 0 & 0 & 0 \\ 0 & 0 & 8.5 & 3.7 & 0 \\ 0 & 6 & 2.7 & 0 & 0 \\ 0 & 4.6 & 0 & 4.8 & 9.4 \\ 0 & 0 & 5.6 & 0 & 1 \end{bmatrix}$$

$$V = [1 \quad 4.6 \quad 6 \quad 4.6 \quad 8.5 \quad 2.7 \quad 5.6 \quad 3.7 \quad 4.8 \quad 9.4 \quad 1]^T$$

$$IA = [0 \quad 0 \quad 2 \quad 3 \quad 1 \quad 2 \quad 4 \quad 1 \quad 3 \quad 3 \quad 4]^T$$

$$JA = [0 \quad 1 \quad 1 \quad 1 \quad 2 \quad 2 \quad 2 \quad 3 \quad 3 \quad 4 \quad 4]^T$$

Polls 3 & 4

- 3 - Which size would you generally use to store integers?
 - A) 16-bit
 - B) 32-bit
 - C) 64-bit
- 4 - Which size would you generally use to store floating-point numbers?
 - A) 16-bit
 - B) 32-bit
 - C) 64-bit

Polls 3 & 4

- 3 - Which size would you generally use to store integers?
 - A) 16-bit
 - **B) 32-bit**
 - C) 64-bit
- 4 - Which size would you generally use to store floating-point numbers?
 - A) 16-bit
 - B) 32-bit
 - **C) 64-bit**
- What is the memory cost in bytes for a COO matrix with 16 entries?

Polls 3 & 4

- 3 - Which size would you generally use to store integers?
 - A) 16-bit
 - **B) 32-bit**
 - C) 64-bit
- 4 - Which size would you generally use to store floating-point numbers?
 - A) 16-bit
 - B) 32-bit
 - **C) 64-bit**
- What is the memory cost in bytes for a COO matrix with 16 entries?
 - $16 * (2 * 4 \text{ (ints)} + 8 \text{ (double)}) = 256 \text{ B}$

Storing Sparse Matrices - CCS

- CCS: Compressed Column Storage

- V , IA are the same as in COO
- JA is a vector of size $(n + 1)$
- JA points to the index of IA where the next column starts
- Storage requirement: $2m + n + 1$

$$\begin{bmatrix} 1 & 4.6 & 0 & 0 & 0 \\ 0 & 0 & 8.5 & 3.7 & 0 \\ 0 & 6 & 2.7 & 0 & 0 \\ 0 & 4.6 & 0 & 4.8 & 9.4 \\ 0 & 0 & 5.6 & 0 & 1 \end{bmatrix}$$

$$V = [1 \quad 4.6 \quad 6 \quad 4.6 \quad 8.5 \quad 2.7 \quad 5.6 \quad 3.7 \quad 4.8 \quad 9.4 \quad 1]^T$$

$$IA = [0 \quad 0 \quad 2 \quad 3 \quad 1 \quad 2 \quad 4 \quad 1 \quad 3 \quad 3 \quad 4]^T$$

$$JA = [0 \quad 1 \quad 4 \quad 7 \quad 9 \quad 11]^T$$

Storing Sparse Matrices – Special Cases

- CRS: **Compressed Row Storage**
 - Same as CCS, exchanging row and column (and the roles of IA , JA)
- Eigen:
 - Defaults to column-major storage
 - Variation of CCS
 - Adds a buffer between each column to insert new elements
 - Additional vector to store the number of non-zero entries
 - `SparseMatrix::makeCompressed()` transforms to standard CCS

Classical Iterative Methods

- Iterative Method: given $Ax = b$, generate a sequence $\{x^{(k)}\}$ which converges to $x = A^{-1}b$
 - Without explicitly calculating A^{-1} - Why?
- Instead of solving $Ax = b$, replace by a simpler $Kx_0 = b$
 - x_0 approximates x as $x = x_0 + z$
 - Plugging back in: $A(x_0 + z) = b \rightarrow Az = b - Ax_0$
- Use the simpler K again
 - $Kz_0 = b - Ax_0$ leading to a new approximation $x_1 = x_0 + z_0$
 - In general: $x_{i+1} = x_i + K^{-1}(b - Ax_i)$
- Examples:
 - Jacobi: $K = D$
 - Gauss-Seidel: $K = D + L$

Richardson Iteration

- Even simpler approximation: $K = I$

$$x_{k+1} = x_k + K^{-1}(b - Ax_k)$$

$$x_{k+1} = x_k + b - Ax_k$$

$$x_{k+1} = x_k + r_k$$

- One way of rewriting

$$x_{k+1} = (I - A)x_k + b$$

- Residuals

$$x_{k+1} = x_k + r_k$$

$$\Leftrightarrow r_{k+1} = (I - A)r_k = (I - A)(I - A)r_{k-1}$$

$$\Leftrightarrow r_{k+1} = (I - A)^{k+1}r_0$$

- Convergence when

$$\|I - A\| < 1$$

Richardson Iteration - Subspace

- The total iteration

$$x_{k+1} = x_k + r_k$$

$$x_{k+1} = x_{k-1} + r_{k-1} + r_k$$

$$x_{k+1} = x_0 + r_0 + \cdots + r_k$$

- Finally, plugging r_i

$$x_{k+1} = x_0 + \sum_{i=0}^k (I - A)^i r_0 = x_0 + z$$

- We can assume w.l.o.g. $x_0 = 0$

$$x_{k+1} \in \text{span}\{r_0, Ar_0, \dots, A^k r_0\}$$

Krylov Subspace - Definition

- Definition: given a non-zero vector v and a non-singular square matrix A , the m -dimensional Krylov subspace is:

$$K^m(A; v) := \text{span}(v, Av, A^2v, \dots, A^{m-1}v)$$

- Krylov Subspace Methods try to:
 - Use all information available in K^m
 - Construct a solution according to some “optimality”
 - More than one option – different methods!

Optimality Approaches

Ritz-Galerkin: Construct x_k s.t. the residual $r_k = b - Ax_k$ is orthogonal to the current subspace

$$r_k \perp K^k(A, r_0)$$

Ex: Conjugate Gradients, Lanczos

Minimum norm residual: Find x_k from least squares

$$\min \|b - Ax_k\|_2; x_k \in K^k(A; r_0)$$

Ex: MINRES, GMRES

Petrov-Galerkin: Construct x_k s.t. the residual $r_k = b - Ax_k$ is orthogonal to some other subspace

Ex: $K^k(A^T; s_0) \rightarrow$ Bi-CG

Also: minimum norm error

Conjugate Gradients (CG) - Motivation

- Assume that A is symmetric positive definite:

$$A = A^T; \quad x^T A x > 0 \quad \forall x \neq 0$$

- We can define a function $\phi(x)$

$$\phi(x) = \frac{1}{2} x^T A x - x^T b$$

- Such that:

$$Ax = b \Leftrightarrow \min_x \phi(x)$$

- Algorithm idea: from a given x_k , **construct x_{k+1} from a search direction p_k and optimality criterion α**

$$x_{k+1} = x_k + \alpha p_k$$

- If $p_k = r_k$, we have the **method of steepest descent**

CG – A-norm

- Definition:

$$\|x\|_A = \sqrt{x^t A x}$$

- We can define the error

$$e_k = x_k - x_*$$

- It can be shown that

$$\phi(x_k) = \frac{1}{2} \|e_k\|_A^2 + \phi(x_*)$$

- CG is [Trefethen & Bau]: a system of recurrence formulas generating the unique sequence $x_k \in K^k$ minimizing $\|e_k\|_A$
 - Trick – Ritz-Galerkin: $r_k \perp K^k(A, r_0)$. Therefore it is “conjugate”!
 - Minimizes the A-norm error, not the residual!

CG – Basic Algorithm (Hestenes & Stiefel)

$$x_0 = 0; r_0 = b; p_0 = r_0$$

for $k = 1, 2, \dots$

$$\alpha_k = \frac{r_{k-1}^T r_{k-1}}{p_{k-1}^T A p_{k-1}}$$

1) step length

$$x_k = x_{k-1} + \alpha_k p_{k-1}$$

2) approximate solution

$$r_k = r_{k-1} - \alpha_k A p_{k-1}$$

3) residual

$$\beta_k = \frac{r_k^T r_k}{r_{k-1}^T r_{k-1}}$$

4) improvement this step

$$p_k = r_k + \beta_k p_{k-1}$$

5) search direction

Plus: convergence termination criterion!

Polls 5 & 6

$$x_0 = 0; r_0 = b; p_0 = r_0$$

for $k = 1, 2, \dots$

$$\alpha_k = \frac{r_{k-1}^T r_{k-1}}{p_{k-1}^T A p_{k-1}}$$

1) step length

$$x_k = x_{k-1} + \alpha_k p_{k-1}$$

2) approximate solution

$$r_k = r_{k-1} - \alpha_k A p_{k-1}$$

3) residual

$$\beta_k = \frac{r_k^T r_k}{r_{k-1}^T r_{k-1}}$$

4) improvement this step

$$p_k = r_k + \beta_k p_{k-1}$$

5) search direction

- Poll 5: How many matrix-vector products per CG iteration?
- Poll 6: How many vector-vector products per CG iteration?

Polls 5 & 6

$$x_0 = 0; r_0 = b; p_0 = r_0$$

for $k = 1, 2, \dots$

$$\alpha_k = \frac{r_{k-1}^T r_{k-1}}{p_{k-1}^T A p_{k-1}}$$

1) step length

$$x_k = x_{k-1} + \alpha_k p_{k-1}$$

2) approximate solution

$$r_k = r_{k-1} - \alpha_k A p_{k-1}$$

3) residual

$$\beta_k = \frac{r_k^T r_k}{r_{k-1}^T r_{k-1}}$$

4) improvement this step

$$p_k = r_k + \beta_k p_{k-1}$$

5) search direction

- Poll 5: How many matrix-vector products per CG iteration? **1**
- Poll 6: How many vector-vector products per CG iteration? **3**

CG – Strengths and Weaknesses

- Strengths
 - Only one matrix-vector operation per iteration
 - Simple implementation
 - **Very fast convergence** (if the eigenvalues are well distributed...)
 - Low memory requirements
- Weaknesses
 - **Only symmetric positive definite (s.p.d.) matrices**

Krylov Subspace - Definition

- Definition: given a non-zero vector v and a non-singular square matrix A , the m -dimensional Krylov subspace is:

$$K^m(A; v) := \text{span}(v, Av, A^2v, \dots, A^{m-1}v)$$

- Krylov Subspace Methods try to:
 - Use all information available in K^m
 - Construct a solution according to some “optimality”
 - More than one option – different methods!

Bases on Krylov Subspace – Arnoldi Iteration

- Problem with $A^k r_0$ -> becomes almost linearly dependent
- **Arnoldi algorithm** (Modified Gram-Schmidt):

$$v_1 = r_0 / \|r_0\|_2$$

first Krylov vector

for $j = 1, \dots, m - 1$

new candidate vector

$$t = Av_j$$

for $i = 1, \dots, j$

$$h_{i,j} = v_i^T t$$

subtract the projections on previous vectors

$$t = t - h_{i,j} v_i$$

$$h_{j+1,j} = \|t\|_2$$

$$v_{j+1} = t / h_{j+1,j}$$

- In practice: Householder reflections

Bases on Krylov Subspace – Hessenberg Matrices

- We just generated an $m \times m$ matrix H_m such that:

$$V_m^T A V_m = H_m$$

- H_m is upper Hessenberg

$$H_m = \begin{bmatrix} \blacksquare & \blacksquare & \blacksquare & \blacksquare & \blacksquare \\ \blacksquare & \blacksquare & \blacksquare & \blacksquare & \blacksquare \\ 0 & \blacksquare & \blacksquare & \blacksquare & \blacksquare \\ 0 & 0 & \blacksquare & \blacksquare & \blacksquare \\ 0 & 0 & 0 & \blacksquare & \blacksquare \end{bmatrix}$$

- Also note:

$$A V_{m-1} = V_m \bar{H}_m$$


where \bar{H}_m has an additional row with only $h_{m+1,m}$

Generalized Minimal Residual Method (GMRES)

- Objective: $\min \|b - Ax_k\|_2; x_k \in K^k(A; r_0)$
- Approach:
 - 1 - Generate Basis $\langle \text{Arnoldi iteration} \rangle$
 - 2 - Optimality Constraint $\min_y \|\beta e_1 - \bar{H}_m y\|_2; \beta = \|r_0\|_2$
 - 3 - Solution construction $x_m = x_0 + V_m y$
- Key insight of GMRES: introduce the **optimality constraint into the basis generation**
 - Usually implemented by Givens rotations during the base generation
- Restarted GMRES: GMRES(m)
 - Only perform GMRES up to dimension m, then restart with current solution x_m

GMRES – Convergence


- Matrix with uniform real spectrum



Source: van der Vorst: Iterative Krylov Methods for Large Linear Systems, 1st ed., 2003

GMRES – Convergence

- Defective (non-diagonalizable) matrix



Source: van der Vorst: Iterative Krylov Methods for Large Linear Systems, 1st ed., 2003

GMRES – Strengths and Weaknesses

- Strengths
 - The residual is **non-increasing**
 - Only “good” breakdowns (if $x_{j < m}$ is already the exact solution)
 - Exact for $m = n$
 - In practice, it is the **most robust method**
- Weaknesses
 - **Memory requirement** (dense in m)
 - Also not the fastest
 - Not in Eigen! (although: Eigen-unsupported)
 - Very hard to prove convergence in arbitrary cases

Bi-Conjugate Gradients (Bi-CG)

- Is it possible to avoid building the m dimensional Hessenberg matrix?
 - We would like a simple recursion like CG, but for non-s.p.d. matrices!
- Insight: build a similar scheme to CG, but searching now on the subspace $K^k(A^T; s_0)$
 - Petrov-Galerkin
- Disadvantages:
 - Requires computing $A^T x$
 - Subject to **serious breakdown**: if the dot product between one vector in $K^k(A^T; s_0)$ and other in $K^k(A; r_0)$ (where the solution is being built) is zero, the method fails.
 - Convergence can be erratic

BiCG – Basic Algorithm

$x_0 = 0; p_0 = r_0 = b; q_0 = s_0 = \text{arbitrary}$

for $k = 1, 2, \dots$

$$\alpha_k = \frac{s_{k-1}^T r_{k-1}}{q_{k-1}^T A p_{k-1}} \quad \leftarrow \text{serious breakdown}$$

$$x_k = x_{k-1} + \alpha_k p_{k-1}$$

$$r_k = r_{k-1} - \alpha_k A p_{k-1}$$

$$s_k = s_{k-1} - \alpha_k A^T q_{k-1}$$

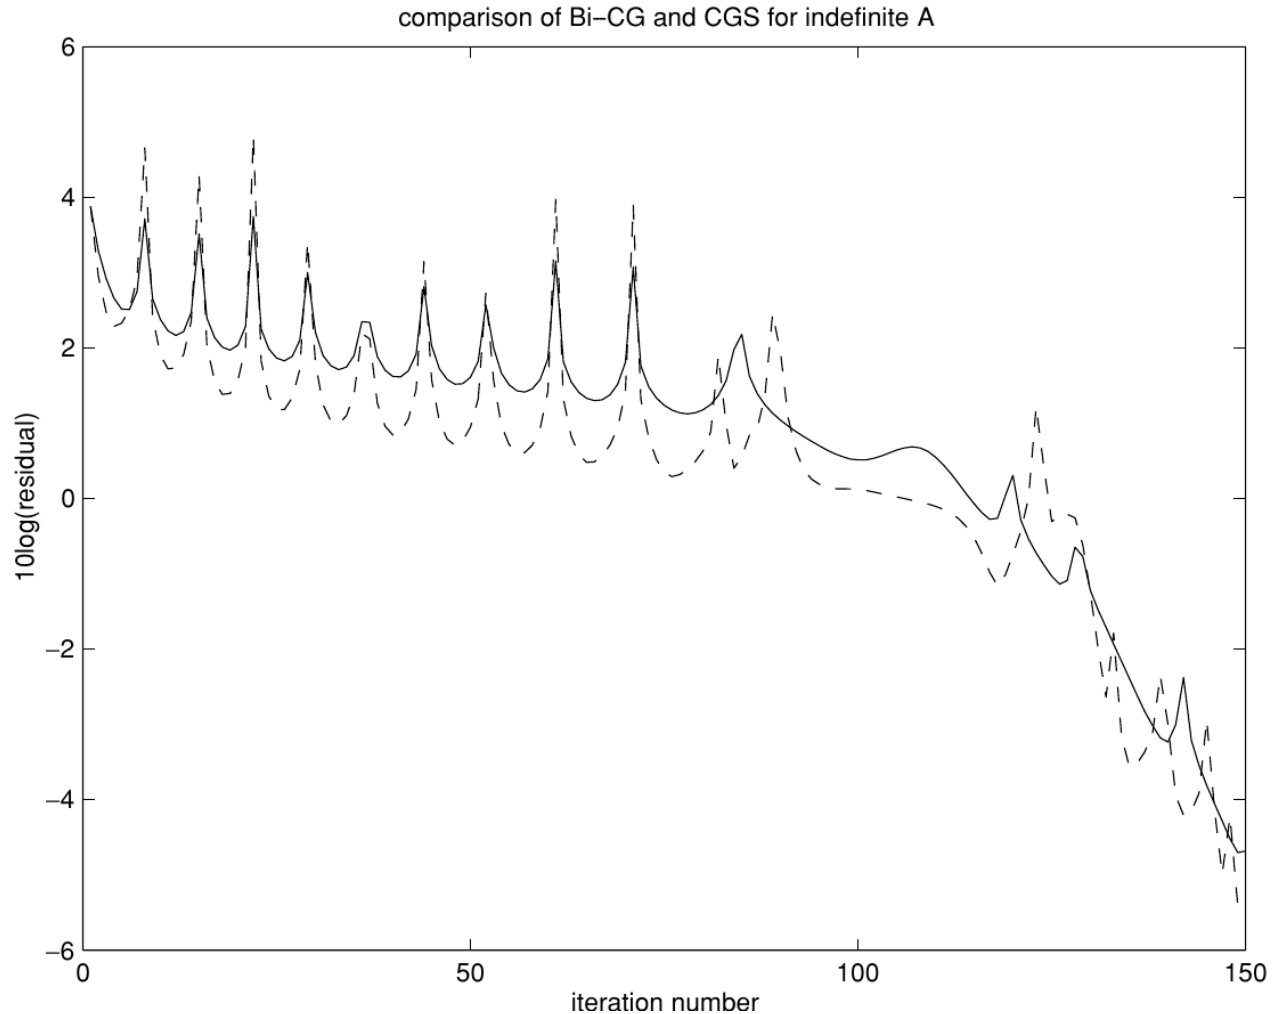
$$\beta_k = \frac{s_k^T r_k}{s_{k-1}^T r_{k-1}}$$

$$p_k = r_k + \beta_k p_{k-1}$$

$$q_k = s_k + \beta_k q_{k-1}$$

BiCG – Erratic Convergence

- Indefinite matrix



Source: van der Vorst: Iterative Krylov Methods for Large Linear Systems, 1st ed., 2003

Bi-CG stabilized (Bi-CGSTAB)

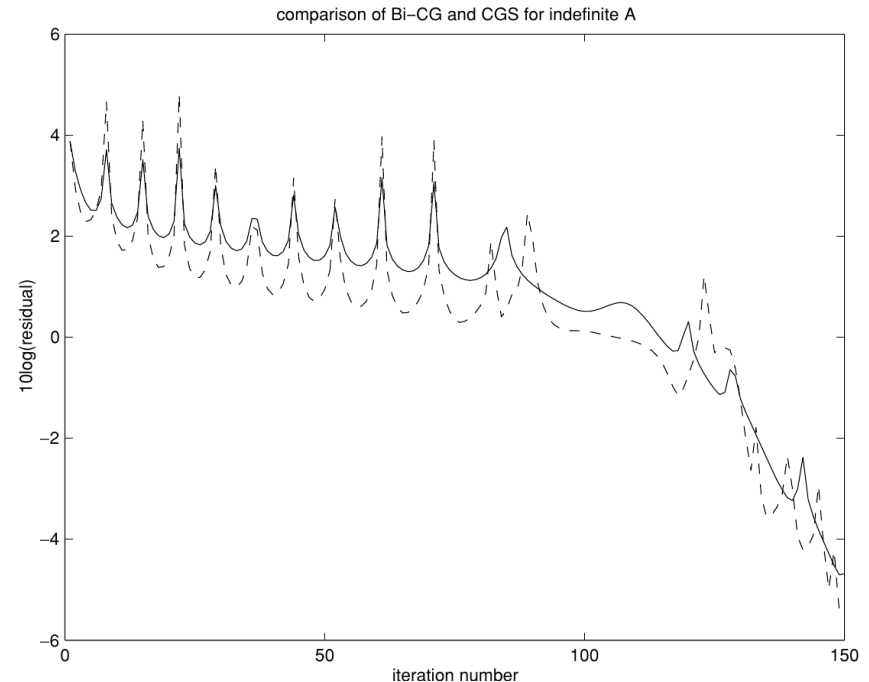

- Problems Bi-CGSTAB addresses:
 - Avoids computing $A^T x$
 - “Smoothens” convergence
- Intuition: combine a **GMRES(1) step after each Bi-CG** step
 - Natural extension: Bi-CGSTAB(l) uses GMRES(l)
- Still susceptible to serious breakdown!
- We will stop at that, the method is (even more) complicated...
 - But: available on EIGEN!

Quiz – Question 5

- What is the benefit of preconditioning a problem before solving it?

Preconditioners - Motivation

- We have seen problematic convergence behavior depending on the **spectrum** of A



- Rule of thumb: iterative methods perform better when the **eigenvalues are clustered**

Preconditioners - Definition

- Construct a matrix K such that the system

$$K^{-1}Ax = K^{-1}b$$

is more easily solvable $\rightarrow K^{-1}A$ has better spectral properties


- Note

- K is in some sense an approximation to A
- K^{-1} is simple to compute: We do not want to store $K^{-1}A$!
- $Ky = z$ is simpler to solve
- Also: the Jacobi method is a preconditioned Richardson iteration

$$D^{-1}Ax = D^{-1}b$$

Preconditioners - Example


- Diagonal scaling/Jacobi preconditioner: $K = \text{diag}(A)$
 - Applied to the CG method



Source: Trefethen, Bau: Numerical Linear Algebra, 1st ed., 1997

Poll 7

- Are both methods converging?



Source: Trefethen, Bau: Numerical Linear Algebra, 1st ed., 1997

Preconditioners - Types

- Left Preconditioning

$$K^{-1}Ax = K^{-1}b$$

- Note: this means that **the residual changes**: $\bar{r}_0 = K^{-1}(b - Ax_0)$

- Right Preconditioning

$$AK^{-1}Kx = b$$

- Note: now the solution lies in a different Krylov subspace

- Mixed/Two-sided Preconditioning

$$K_1^{-1}AK_2^{-1}K_2^{-1}x = K_1^{-1}b$$

Incomplete LU factorization

- We know that, given an LU factorization, it is easy to solve

$$LUx = b$$

- What if we calculate a simpler, **incomplete LU**?

$$A \approx LU$$

- Simplest idea: only calculate L and U for **elements with the same indices that are non-zero in A**
 - Keep the same sparsity
 - Does not always help convergence...
- More advanced schemes: set **tolerances to drop elements**
 - Depends on the heuristics...

Note – Direct solvers are still relevant!

The Sparse LU Challenge

- Given a square A , find the permutation matrices P and Q such that the factorization $PAQ^T = LU$ is:
 - Reasonably stable
 - L and U are close to being optimally sparse
- Many approaches
 - Example: http://eigen.tuxfamily.org/dox/classEigen_1_1SparseLU.html
- Can benefit from BLAS!
- Some sparsity will be lost...

Take-home message – version 2

- Krylov methods can be useful when there are **large, sparse** matrices
 - We do not want to lose sparsity
 - We do not want to pay the $O(n^3)$ price
 - *But not always better than direct methods...*
- Convergence can be hard
 - Choosing the **correct method for the problem**
 - Equal parts “art” and “science” – experience & experimentation are key!
 - *Some idea about the spectrum of the matrix is helpful*
- Preconditioners can significantly increase convergence rate
 - Also make your problem more stable
 - *This is where a lot of mathemagic is hidden*

Rules of Thumb for Choosing Methods

- If the matrix is symmetric positive definite:
 - Conjugate Gradients
- If the problem is “nasty”:
 - GMRES + lots of memory + lots of computing time
- If performance is key:
 - Bi-CGSTAB (or other hybrid methods) + lots of robustness testing
- **Always investigate preconditioners**
 - But remember that they involve trial-and-error
- If you can afford the memory, test sparse direct methods

Quiz

- Q1: Represent the following matrix in the compressed row storage (CRS) format

$$\begin{bmatrix} 1 & 4.6 & 0 & 0 & 0 \\ 0 & 0 & 8.5 & 3.7 & 0 \\ 0 & 6 & 2.7 & 0 & 0 \\ 0 & 4.6 & 0 & 4.8 & 9.4 \\ 0 & 0 & 5.6 & 0 & 1 \end{bmatrix}$$

- Q2: Under which conditions would you use the CG method instead of GMRES?
- Q3: What information from a matrix A could be useful to help choosing an adequate solver?
- Q4: Is it possible to generate true random numbers from a digital computer (e.g. x86 architecture)?
- Q5: How would you estimate the value of π using random numbers?

Next stop

- Random Number Generation
- Introduction to Monte Carlo