

# Numerical Simulation and Scientific Computing I

## Lecture 7: Random Numbers & Monte Carlo Integration



Clemens Etl, Paul Manstetten,  
and Josef Weinbub



Institute for Microelectronics  
TU Wien

[nssc@ime.tuwien.ac.at](mailto:nssc@ime.tuwien.ac.at)

# Quiz

- Q1: Represent the following matrix in the CRS format

$$\begin{bmatrix} 1 & 4.6 & 0 & 0 & 0 \\ 0 & 0 & 8.5 & 3.7 & 0 \\ 0 & 6 & 2.7 & 0 & 0 \\ 0 & 4.6 & 0 & 4.8 & 9.4 \\ 0 & 0 & 5.6 & 0 & 1 \end{bmatrix}$$

# Quiz

- Q1: Represent the following matrix in the CRS format

$$\begin{bmatrix} 1 & 4.6 & 0 & 0 & 0 \\ 0 & 0 & 8.5 & 3.7 & 0 \\ 0 & 6 & 2.7 & 0 & 0 \\ 0 & 4.6 & 0 & 4.8 & 9.4 \\ 0 & 0 & 5.6 & 0 & 1 \end{bmatrix}$$

- 1<sup>st</sup> step: Put the values in row-major order

$$V = [1 \quad 4.6 \quad 8.5 \quad 3.7 \quad 6 \quad 2.7 \quad 4.6 \quad 4.8 \quad 9.4 \quad 5.6 \quad 1]^T$$

# Quiz

- Q1: Represent the following matrix in the CRS format

$$\begin{bmatrix} 1 & 4.6 & 0 & 0 & 0 \\ 0 & 0 & 8.5 & 3.7 & 0 \\ 0 & 6 & 2.7 & 0 & 0 \\ 0 & 4.6 & 0 & 4.8 & 9.4 \\ 0 & 0 & 5.6 & 0 & 1 \end{bmatrix}$$

- 2<sup>nd</sup> step: Assign the corresponding column indices

$$V = [1 \quad 4.6 \quad 8.5 \quad 3.7 \quad 6 \quad 2.7 \quad 4.6 \quad 4.8 \quad 9.4 \quad 5.6 \quad 1]^T$$

$$JA = [0 \quad 1 \quad 2 \quad 3 \quad 1 \quad 2 \quad 1 \quad 3 \quad 4 \quad 2 \quad 4]^T$$

# Quiz

- Q1: Represent the following matrix in the CRS format

$$\begin{bmatrix} 1 & 4.6 & 0 & 0 & 0 \\ 0 & 0 & 8.5 & 3.7 & 0 \\ 0 & 6 & 2.7 & 0 & 0 \\ 0 & 4.6 & 0 & 4.8 & 9.4 \\ 0 & 0 & 5.6 & 0 & 1 \end{bmatrix}$$

- 3<sup>rd</sup> step: Identify the indices in  $V$  where a new row starts

$$V = [1 \quad 4.6 \quad 8.5 \quad 3.7 \quad 6 \quad 2.7 \quad 4.6 \quad 4.8 \quad 9.4 \quad 5.6 \quad 1]^T$$

$$JA = [0 \quad 1 \quad 2 \quad 3 \quad 1 \quad 2 \quad 1 \quad 3 \quad 4 \quad 2 \quad 4]^T$$

$$IA = [0 \quad 2 \quad 4 \quad 6 \quad 9 \quad 11]^T$$

# Quiz – Poll 1

- Q2: Under which conditions would you use the CG method instead of GMRES?
- A) The matrix is indefinite
- B) The problem benefits from preconditioning
- C) The problem is ill-posed
- D) All eigenvalues are positive
- E) The LU factorization is known

# Quiz – Poll 1

- Q2: Under which conditions would you use the CG method instead of GMRES?
- A) The matrix is indefinite
- B) The problem benefits from preconditioning
- C) The problem is ill-posed
- **D) All eigenvalues are positive**
- E) The LU factorization is known

## Quiz – Poll 2

- Q3: What information from a matrix  $A$  could be useful to help choosing an adequate solver?
- A) Maximum norm
- B) Set of eigenvalues
- C) QR decomposition
- D) Transpose



## Quiz – Poll 2

- Q3: What information from a matrix  $A$  could be useful to help choosing an adequate solver?
- A) Maximum norm
- **B) Set of eigenvalues**
- C) QR decomposition
- D) Transpose

# Outline

- Introduction to Randomness
- Random Number Generators (RNGs)
  - Linear Congruential Generators
  - Mersenne Twister
  - Hardware
- C++11 `<random>`
- Monte Carlo Integration

# Take-home message

---

- Be careful with your choice of RNG
- Seeding is a consequence of pseudo-RNG
- Watch out for pitfalls on parallel systems

# Main References

---

- Random Numbers and Computers
  - Author: Ronald T. Kneusel
  - eBook available:  
[https://catalogplus.tuwien.ac.at:443/UTW:UTW:TN\\_springer\\_s978-3-319-77697-2\\_453445](https://catalogplus.tuwien.ac.at:443/UTW:UTW:TN_springer_s978-3-319-77697-2_453445)

# Additional References

- Handbook of Monte Carlo Methods
  - Authors: D. P. Kroese, T. Taimre, Z. I. Botev
  - eBook available:  
[https://catalogplus.tuwien.ac.at:443/UTW:UTW:TN\\_scopus2-s2.0-84949783693](https://catalogplus.tuwien.ac.at:443/UTW:UTW:TN_scopus2-s2.0-84949783693)
- Monte Carlo Strategies in Scientific Computing
  - Author: Jun S. Liu
  - [https://catalogplus.tuwien.ac.at:443/UTW:UTW:UTW\\_alma2146615720003336](https://catalogplus.tuwien.ac.at:443/UTW:UTW:UTW_alma2146615720003336)

# Randomness


- How to formally define randomness?
- Philosophical conundrum: does randomness even exist?
  - Good thing that it is not our job to figure it out!
- Too hard - operational definition of a random sequence [Kneusel]:
  - a sequence of numbers  $n$ , within some bounded range, where it is not possible to predict  $n_{k+1}$  from any combination of preceding values  $n_i, i = 0, 1, \dots, k$ .

# “True” Random vs. Pseudorandom

- “True” Randomness
  - From here on, we will **assume that these exist**
  - Found in the physical world
  - Examples?
- Pseudorandomness
  - [Kneusel]: A pseudorandom sequence is a **deterministically** generated sequence of numbers that is **indistinguishable** from a true random sequence of numbers
  - Perfect indistinguishability is impossible – but how much is enough?
  - Determinism: introduces the concept of a **seed**
    - **Seed**: starting point of the deterministic sequence

# Uniform Random Numbers

- Even if  $n_{k+1}$  is truly non-predictable, it can still have a distribution – shows up in a histogram
  - Drawing from a **parent distribution**
- Uniform (real) Distribution:  $n_k \in [0,1)$ 
  - Unless otherwise stated, it is **often the default**





# Linear Congruential Generators

- Nomenclature

- $a$ : multiplier
- $c$ : increment
- $m$ : modulus
- $x_0$ : seed

$$x_{n+1} = (ax_n + c) \bmod m$$

- Notes

- Generates integers  $\in [0, m)$ . Floats:  $f_n = \frac{x_n}{m}$
- **Periodicity** is a big problem – can be lower than  $m$
- If  $x_n = x_m$  then  $x_{n+1} = x_{m+1}$
- The method is defined by the choice of  $a$ ,  $c$  and  $m$

- Examples

- GCC:  $a = 1103515245$ ,  $c = 12345$ ,  $m = 2^{31} - 1$
- C++11 minstd\_rand:  $a = 48271$ ,  $c = 0$ ,  $m = 2^{31} - 1$

# LCGs – Poll 3

- What is the issue of using this generator with  $m = 2^{31} - 1$  to generate a double?

$$x_{n+1} = (ax_n + c) \bmod m$$

- A) The periodicity is not high enough to cover the 53-bit significand
- B) It is impossible to seed this generator with a double
- C) You cannot generate a double from an integer
- D) The equation  $f_n = \frac{x_n}{m}$  doesn't hold anymore

## LCGs – Poll 3

- What is the issue of using this generator with  $m = 2^{31} - 1$  to generate a double?

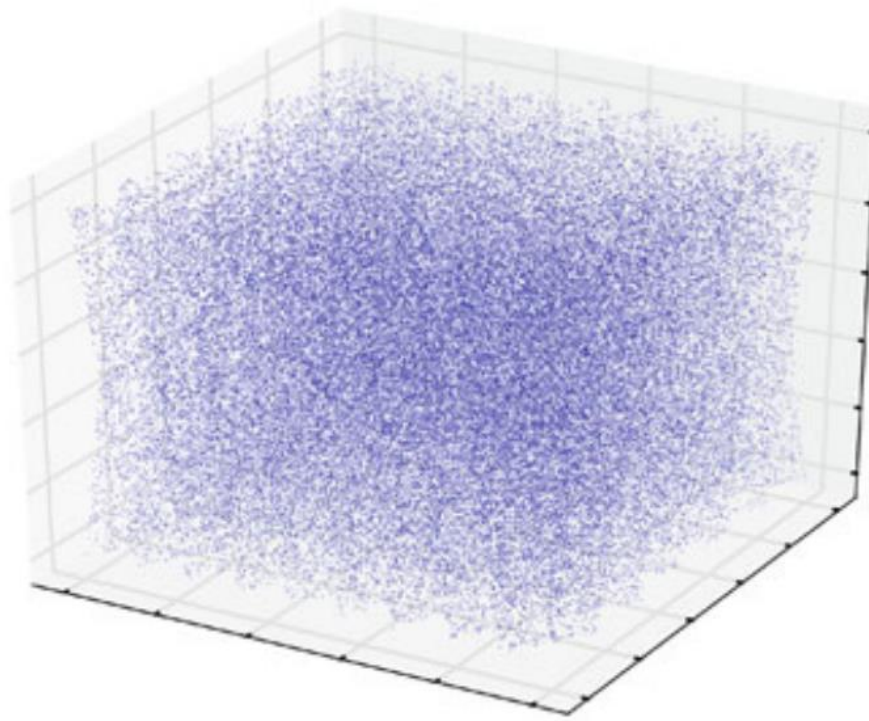
$$x_{n+1} = (ax_n + c) \bmod m$$

- **A) The periodicity is not high enough to cover the 53-bit significand**
- B) It is impossible to seed this generator with a double
- C) You cannot generate a double from an integer
- D) The equation  $f_n = \frac{x_n}{m}$  doesn't hold anymore

# A Lesson from History: RANDU

- How bad can it be if we choose poorly?
  - RANDU:  $a = 65593$ ,  $c = 0$ ,  $m = 2^{31}$

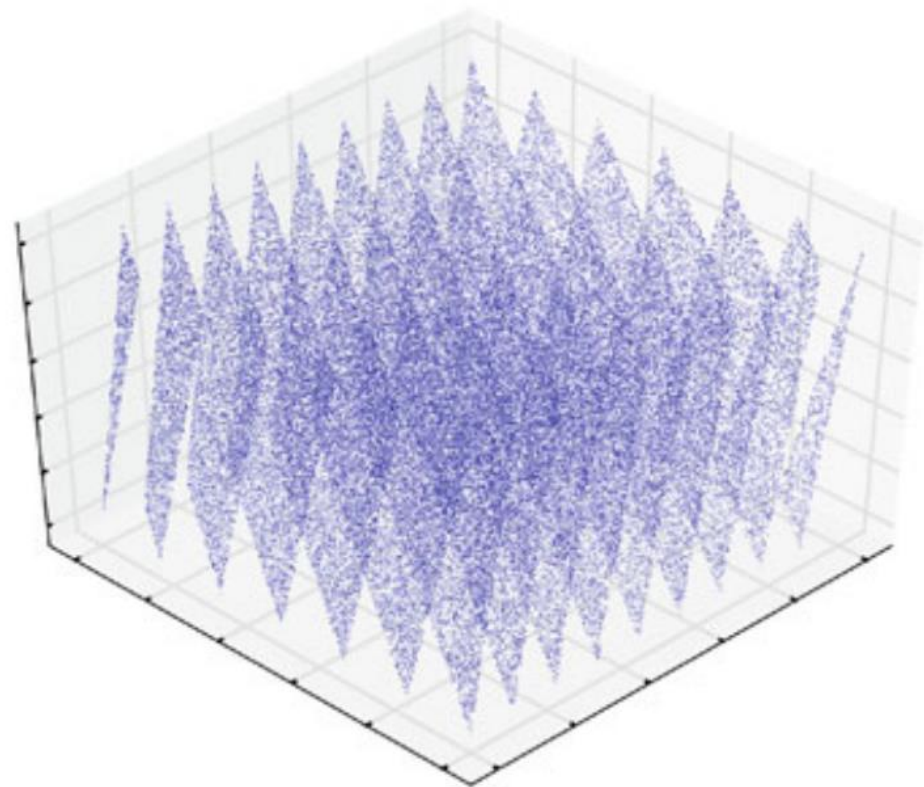

$$x_{n+2} = 6x_{n+1} - 9x_n$$



# A Lesson from History: RANDU


- How bad can it be if we choose poorly?
  - RANDU:  $a = 65593$ ,  $c = 0$ ,  $m = 2^{31}$

$$x_{n+2} = 6x_{n+1} - 9x_n$$



# Mersenne Twister

- MT is the current *de facto standard* for non-cryptographic applications



- Larger memory and construction requirements!

Source: Kneusel: Random Numbers and Computers, 1<sup>st</sup> ed., 2018, Springer

# Choosing your RNG

- Performance is **only part of the equation**

- MINSTD is only 32 bits

- $a$ ,  $c$  and  $m$  are linked
  - **Don't mess around with constants!**

- The performance penalty of MT can be mild (up to 50%)

- Other RNG may be available!

- May involve other performance tradeoffs...

- C rand()

- NOOOOOOOO ( $RAND\_MAX \geq 32767$ )

Generator	Time (s)
MINSTD	$9.164 \pm 0.006$
xorshift32	$11.097 \pm 0.004$
xorshift128+	$11.853 \pm 0.028$
CMWC	$12.671 \pm 0.003$
Middle Weyl	$12.849 \pm 0.028$
xorshift1024*	$12.985 \pm 0.004$
Mersenne Twister	$14.353 \pm 0.007$
KISS64	$17.252 \pm 0.015$
Philox	$18.097 \pm 0.048$
Threefry	$32.442 \pm 0.010$

Source: Kneusel: Random Numbers and Computers, 1<sup>st</sup> ed., 2018, Springer

<https://channel9.msdn.com/Events/GoingNative/2013/rand-Considered-Harmful>

# Choosing your RNG – Poll 4

- How does MT compare wrt. MINSTD?
  - Higher performance: less time to generate 1 RN
- A) Higher performance, higher periodicity, higher memory requirements
- B) Lower performance, higher periodicity, higher memory requirements
- C) Lower performance, lower periodicity, higher memory requirements
- D) Lower performance, higher periodicity, lower memory requirements
- E) Higher performance, lower periodicity, lower memory requirements



# Choosing your RNG – Poll 4

- How does MT compare wrt. MINSTD?
  - Higher performance: less time to generate 1 RN
- A) Higher performance, higher periodicity, higher memory requirements
- **B) Lower performance, higher periodicity, higher memory requirements**
- C) Lower performance, lower periodicity, higher memory requirements
- D) Lower performance, higher periodicity, lower memory requirements
- E) Higher performance, lower periodicity, lower memory requirements

# Choosing your RNG – Poll 5

- For a certain random sequence, in the past we drew the number 189 immediately after the number 25. Now in the same sequence, we just drew the number 25 again.
- A) MT: the next number will be 189. LCG: the next number will be 189
- B) MT: the next number will be random. LCG: the next number will be 189
- C) MT: the next number will be 189. LCG: the next number will be random
- D) MT: the next number will be random. LCG: the next number will be random

# Choosing your RNG – Poll 5

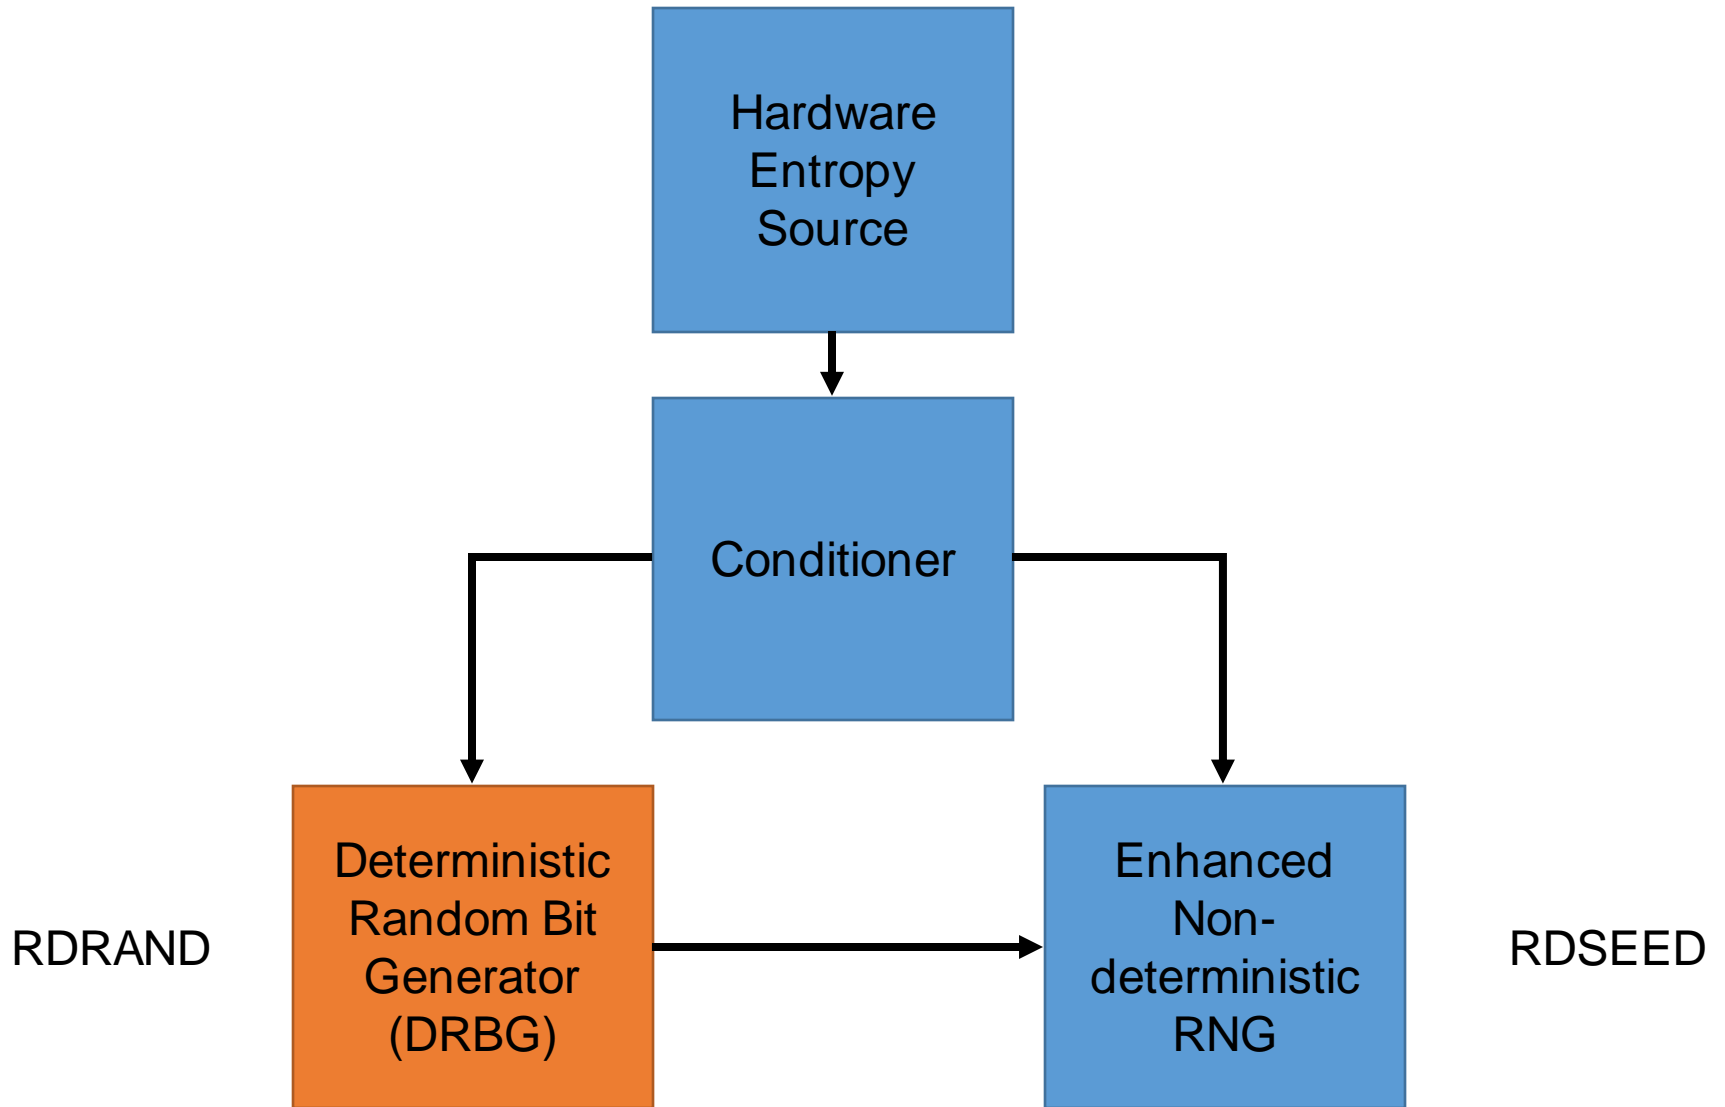
- For a certain random sequence, in the past we drew the number 189 immediately after the number 25. Now in the same sequence, we just drew the number 25 again.
- A) MT: the next number will be 189. LCG: the next number will be 189
- **B) MT: the next number will be random. LCG: the next number will be 189**
- C) MT: the next number will be 189. LCG: the next number will be random
- D) MT: the next number will be random. LCG: the next number will be random

## Quiz Q4 – Poll 6

---

- Is it possible to generate true random numbers from a digital computer (e.g. x86 architecture)?
  - Yes
  - No

# Intel DRNG



# Entropy Pool

- Entropy pool: where you **store the true randomness**
  - E.g.: hardware, /dev/random
- Depletion:
  - Every time you use the pool to generate a RN, you gain information about it
  - The next query might not be random!
  - The pool usually “knows” it is depleted -> needs **more randomness** (time) to replenish

# RDRAND - Example

```
#include <iostream>

int main() {
    unsigned long long int random;



    for (int i = 0; i < 10; ++i) {
        if (__builtin_ia32_rdrand64_step(&random)) break;
    }

    std::cout << random << std::endl;
}
```

- GCC specific: compilation with `g++ -mrdrnd`
- Available in Intel Ivy Bridge and AMD Zen onwards
- Downside: lack of seeding means **losing reproducibility**

# Generating Other Distributions

- The basic algorithms generate a uniform  $u \in [0,1)$
- What if we need **other distributions**?
  - Common example: **normal (Gaussian) distribution**  $N(\mu, \sigma)$
  - $N(\mu, \sigma) = \sigma * N(0,1) + \mu$
- We would like a map  $u \rightarrow z \in N(0,1)$
- Very **common trick**: map  $u_1, u_2 \in [0,1) \rightarrow z_1, z_2 \in N(0,1)$





# Normal Distribution - Box-Muller Algorithm

- Draw 2 uniform samples  $u_1$  and  $u_2$

$$z_1 = \sqrt{-2\ln(u_1)} \cos(2\pi u_2)$$

$$z_2 = \sqrt{-2\ln(u_1)} \sin(2\pi u_2)$$

```
import matplotlib.pyplot as plt
import numpy as np
```


```
u1 = np.random.random(10000)
u2 = np.random.random(10000)
```

```
z1 = np.sqrt(-2*np.log(u1))*np.cos(2*np.pi*u2)
z2 = np.sqrt(-2*np.log(u1))*np.sin(2*np.pi*u2)
```

```
fig, ax = plt.subplots()
ax.hist(z1, 25, density=1, label="z1")
ax.hist(z2, 25, density=1, label="z2")
ax.legend()
```

```
plt.savefig('boxmuller.png', dpi=300)
```

# Normal Distribution - Box-Muller Algorithm



# Random Numbers in C++11

- C++11 introduced the `<random>` header
  - <https://www.youtube.com/watch?v=6DPkyvkMkk8>
- Random number engines and engine adaptors
  - E.g. `std::mersenne_twister_engine`
  - These folks mess around with the constants and the internal state of the RNGs
  - Don't mess around with the given constants
- Predefined RNGs
  - E.g. `std::mt19937_64` is an instantiation of the engine with the correct constants for a 64-bit random number
  - A collection of popular RNGs including both 32 and 64-bit Mersenne Twister

# Random Numbers in C++11

- Random Number Distributions
  - E.g. `std::uniform_int_distribution`
  - It can be imagined as a view of the RNG
  - Given a certain generator, **extract a new random number according to the distribution**
  - Many common choices, **including Gaussian, Poisson...**
- Non-deterministic Random Numbers
  - `std::random_device` generates a non-deterministic uniform 32-bit int
  - It **may** use hardware-specific RNG.
  - Does your implementation actually use RDRAND? `^\\_(ツ)_/`
  - Common application: **generate a good seed for your RNG**
  - It can be used for general purpose, but might become slower

# Random Numbers in C++11 - Example

```
#include <iostream>
#include <random>

int main()
{
    std::random_device rd;
    unsigned int rd_seed = rd();
    std::mt19937_64 gen(rd_seed);
    std::normal_distribution<double> d(0,1);

    for(int i=0; i<5; ++i) {
        std::cout << d(gen) << std::endl;
    }
}
```

# Random Numbers in a Parallel System – Poll 7

- What happens when different RNGs are initialized in a parallel system with the same seed?
- A) Each stream generates different, independent RNs
- B) The result of each stream is random, depending on initialization
- C) All threads read from the same RNG, leading to data races
- D) All streams generate the same sequence of RNs

# Random Numbers in a Parallel System – Poll 7

- What happens when different RNGs are initialized in a parallel system with the same seed?
- A) Each stream generates different, independent RNs
- B) The result of each stream is random, depending on initialization
- C) All threads read from the same RNG, leading to data races
- **D) All streams generate the same sequence of RNs**

# Random Numbers in a Parallel System

- Random number server
  - The program only has a **single RNG** and all threads and processes draw numbers from it
  - Commonly used for GPUs: the CPU generates a list of random numbers which is then loaded to the GPU
- Per-stream generator
  - Assign a **different RNG per thread** or process
  - **Pay attention to the seeds!**
  - One approach: 1 master generator with 1 seed -> 1 seed per process/thread



# Sidenote: CSRNG

- Working definition [Kneusel] of Cryptographically Safe RNG:
  - A CSPRNG is a random number generator that passes the next-bit test and one where an attacker's knowledge of the state of the generator at time  $t$  **makes knowing the state at any previous time impossible**.
- Irrelevant for simulations
- But a requirement for cryptography!
- MT is NOT cryptographically safe

# Rules of Thumb

- If you do not want to think, use the **Mersenne Twister**
  - Consider the 64-bit version when generating doubles...
- Each thread or process needs to have **its own RNG**
  - With **different seed!**
- Avoid LCG if possible
  - Every time you use C-style `rand()`, a transistor cries
  - In general, their period is too short for large scale simulations
- Hardware RNGs might make sense in many applications!
  - Watch out for **entropy pool depletion**
- Correct seeding can be hard – but allows for reproducibility!

# Monte Carlo

- What is/Where is Monte Carlo?
- Term dates to the **Manhattan Project**
- Very loosely defined: any type of **algorithm which uses random sampling to achieve a numerical result**
  - The problem itself might be deterministic!
  - Las Vegas methods (subclass): uses randomness but achieves a deterministic result
- “Monte Carlo” has in practice many definitions for each field


## Quiz – Question 5

---

- How would you estimate the value of  $\pi$  using random numbers?

# Monte Carlo Integration – one way

- Draw a circle inside the unit square
  - $g(x) = \sqrt{1 - x^2}$
- Generate pairs of random numbers
  - $x, y \in [0, 1)$
- Count the number of hits inside
  - $y < g(x)$
- Calculate the ratio:
  - $\frac{\pi}{4} \approx \frac{\#_{inside}}{\#_{total}}$



# Monte Carlo Integration – another way

$$I = \int_D g(\mathbf{x}) d\mathbf{x}$$



$$\hat{I}_m = \frac{1}{m} \{g(\mathbf{x}^{(1)}) + \cdots + g(\mathbf{x}^{(m)})\}$$

- Where  $\mathbf{x}^{(n)}$  are drawn uniformly from  $D$
- For the circle,  $D: x \in [0,1)$  and  $g(x) = \sqrt{1 - x^2}$

# Take-home message – version 2

- Be **careful with your choice** of RNG
  - The Mersenne Twister is the standard for a reason
  - Hardware RNGs are becoming more useful
  - Older LCG are to be avoided whenever possible
- **Seeding** is a consequence of pseudo-RNG
  - Not always a bad thing – reproducibility!
- Watch out for pitfalls on parallel systems
  - Usually, you must make sure that each process/thread is generating a **different random number sequence**

# Quiz

- Q1: Consider a 32-bit LCG. How many random numbers can you generate before the sequence repeats?
- Q2: What are the consequences of entropy pool depletion? How would you handle it?
- Q3: What are the requirements to initialize correctly one RNG per thread?
- Q4: Which STL containers are usually implemented using hashes? What is the advantage of doing so?
- Q5: How are elements accessed on a `std::list`?



# Next stop

- Lectures swapped: OpenMP -> 14.12.2023
- Algorithm Analysis
- Data Structures